

image-classification

April 22, 2023

Sanjana Jadhav (CS 4375.004)

1 Image Classification using Deep Learning

In this notebook, we will try to classify cards using the following three models: Sequential Model, Convolutional Neural Networks Model, & Pre-Trained (Transfer Learning) Model

Dataset: <https://www.kaggle.com/datasets/gpiosenka/cards-image-datasetclassification> (this dataset was already split into train and test)

Other Sources: Slides, Textbook, Professor Mazidi's Github + other sources listed below

1.0.1 Data Processing + Visualization

import numpy, pandas, & tensorflow

- numpy: numerical computing library
- pandas: data analysis library
- tensorflow: used to build/train neural network models

```
[2]: import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt
```

import train/test data

- use tf to import train/test data
- source: https://www.tensorflow.org/api_docs/python/tf/keras/utils/image_dataset_from_directory
- important parameters: image size (224 x 224), seed=1234 (for reproducibility), batch size is 32

```
[4]: train = tf.keras.preprocessing.image_dataset_from_directory(
    "/kaggle/input/cards-image-datasetclassification/train",
    labels="inferred",
    label_mode="categorical",
    color_mode="rgb",
```

```

        image_size=(224, 224),
        seed=1234,
        batch_size=32,
    )

test = tf.keras.preprocessing.image_dataset_from_directory(
    "/kaggle/input/cards-image-datasetclassification/test",
    labels="inferred",
    label_mode="categorical",
    color_mode="rgb",
    image_size=(224, 224),
    seed=1234,
    batch_size=32,
)

```

Found 7624 files belonging to 53 classes.

Found 265 files belonging to 53 classes.

1.0.2 Print/Plot Class Distribution

Sources Used:

- <https://docs.python.org/3/library/os.html>
- <https://realpython.com/python-zip-function/>
- <https://stackoverflow.com/questions/28663856/how-do-i-count-the-occurrence-of-a-certain-item-in-an-ndarray>
- https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.bar.html
- https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.figure.html
- https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.xticks.html

```

[33]: print(train.class_names)

# set number of classes
num_classes = len(train.class_names)

```

```

['ace of clubs', 'ace of diamonds', 'ace of hearts', 'ace of spades', 'eight of
clubs', 'eight of diamonds', 'eight of hearts', 'eight of spades', 'five of
clubs', 'five of diamonds', 'five of hearts', 'five of spades', 'four of clubs',
'four of diamonds', 'four of hearts', 'four of spades', 'jack of clubs', 'jack
of diamonds', 'jack of hearts', 'jack of spades', 'joker', 'king of clubs',
'king of diamonds', 'king of hearts', 'king of spades', 'nine of clubs', 'nine
of diamonds', 'nine of hearts', 'nine of spades', 'queen of clubs', 'queen of
diamonds', 'queen of hearts', 'queen of spades', 'seven of clubs', 'seven of
diamonds', 'seven of hearts', 'seven of spades', 'six of clubs', 'six of
diamonds', 'six of hearts', 'six of spades', 'ten of clubs', 'ten of diamonds',
'ten of hearts', 'ten of spades', 'three of clubs', 'three of diamonds', 'three
of hearts', 'three of spades', 'two of clubs', 'two of diamonds', 'two of
hearts', 'two of spades']

```

```
[5]: import os


# Get the class names and indices
class_names = os.listdir("/kaggle/input/cards-image-datasetclassification/
↳train")
class_indices = dict(zip(class_names, range(len(class_names))))

# Get class counts
class_counts = np.zeros(len(class_names))
for images, labels in train:
    class_counts += np.sum(labels.numpy(), axis=0)

# Increase size of graph for readability
plt.figure(figsize=(20, 6))

# Create a bar graph
plt.bar(class_names, class_counts, width=0.6, align='center', color='deeppink')
plt.title('Distribution of Target Classes')
plt.xlabel('Target Class')
plt.ylabel('Count')
plt.xticks(rotation=90)

# Print graph
plt.show()
```



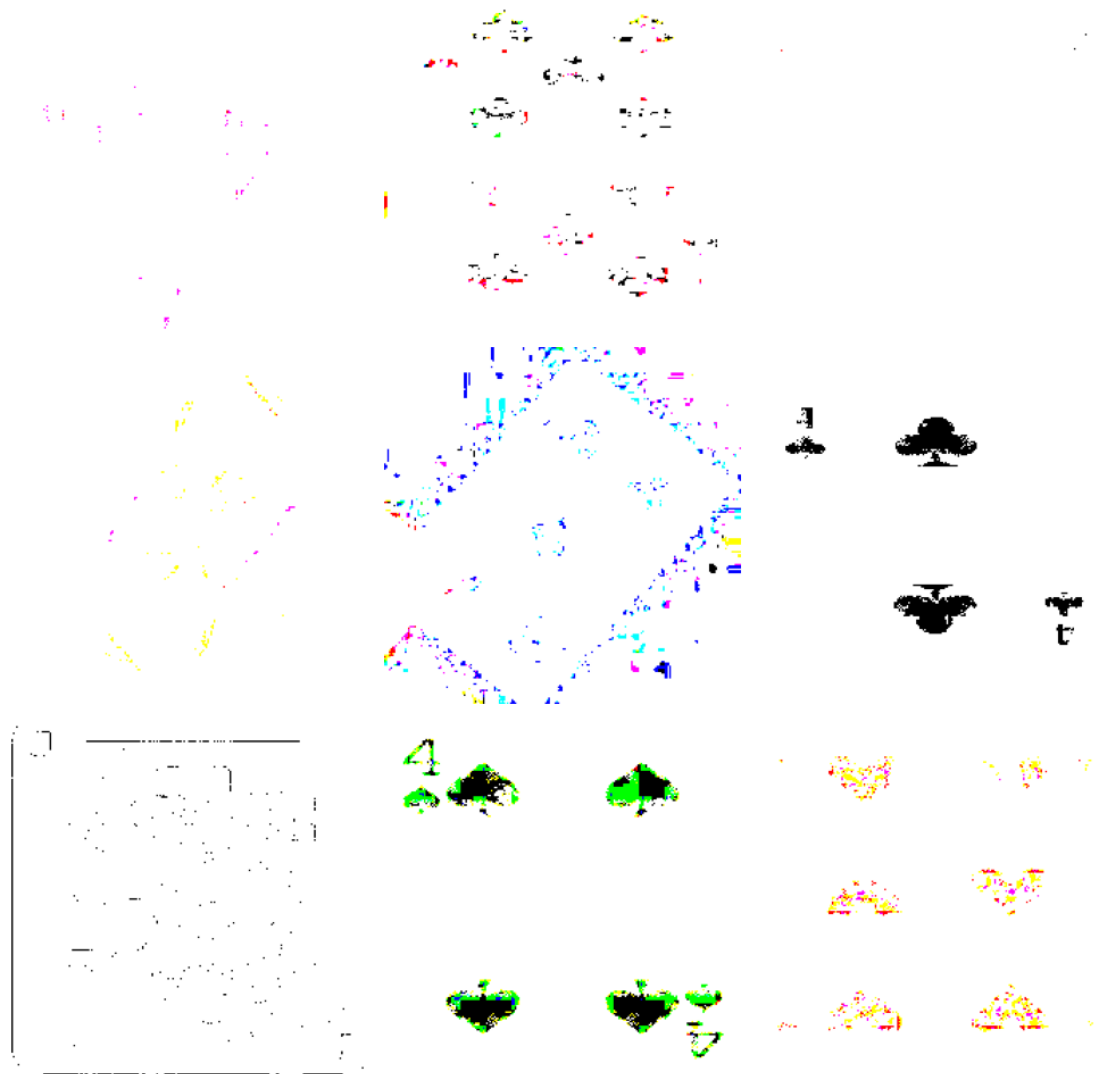
1.0.3 Plot Sample Images from Training Data

Source Used:

- <https://github.com/kjmazidi/Intro-to-Deep-Learning/blob/main/Part%202%20Deep%20Learning%20with%20Cats%20and%20dogs%20classification.ipynb>

- As we can see, it is hard to see the images.

```
[27]: def plotImages(images_arr):  
    fig, axes = plt.subplots(3, 3, figsize=(10,10))  
    axes = axes.flatten()  
    for img, ax in zip(images_arr, axes):  
        ax.imshow(img)  
        ax.axis('off')  
    plt.tight_layout()  
    plt.show()  
  
    # Get a batch of images and labels from the dataset  
    images, labels = next(iter(train))  
    image_batch = images.numpy()  
  
    # Plot the images in a 3x3 grid  
    plotImages(image_batch[:9])
```



1.0.4 Sequential Model

- input layer: image size is 224 x 224 and channels = 3
- rescale the image between [0,1]
- flattens tensor into 1d array
- create a dense layer with 256 neurons and relu activation
- drop (0.5) to prevent overfitting
- create another dense layer with 128 neurons and relu activation
- drop (0.5) to prevent overfitting
- create another a dense layer with 53 neurons (num_classes) and softmax activation (used for multi-class classification to output probabilities)

```
[34]: model = tf.keras.Sequential([
    tf.keras.layers.InputLayer(input_shape=(224, 224, 3)),
    tf.keras.layers.Rescaling(1./255),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(256, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])
```

Compile/Train Model

```
[35]: # Compile Model
# Loss: categorical_crossentropy (used for multi-class classification for
#      ↪ labels that must fit in one class)
# Optimizer: rmsprop
# Metrics: Accuracy
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])

# Train model
# Train model using the training data, 10 epochs, a batch size of 32, and use
#      ↪ the test data
history = model.fit(train,
                    epochs=10,
                    batch_size=32,
                    validation_data=test)
```

Epoch 1/10

239/239 [=====] - 98s 406ms/step - loss: 16.4403 - accuracy: 0.0197 - val_loss: 3.9690 - val_accuracy: 0.0189

Epoch 2/10

239/239 [=====] - 97s 407ms/step - loss: 4.0418 - accuracy: 0.0219 - val_loss: 3.9703 - val_accuracy: 0.0189

Epoch 3/10

239/239 [=====] - 98s 410ms/step - loss: 3.9663 - accuracy: 0.0233 - val_loss: 3.9647 - val_accuracy: 0.0226

Epoch 4/10

239/239 [=====] - 99s 412ms/step - loss: 3.9661 - accuracy: 0.0239 - val_loss: 3.9701 - val_accuracy: 0.0189

Epoch 5/10

239/239 [=====] - 96s 402ms/step - loss: 3.9651 - accuracy: 0.0237 - val_loss: 3.9686 - val_accuracy: 0.0189

Epoch 6/10

239/239 [=====] - 95s 398ms/step - loss: 3.9636 -

```

accuracy: 0.0240 - val_loss: 4.0178 - val_accuracy: 0.0189
Epoch 7/10
239/239 [=====] - 95s 397ms/step - loss: 4.1082 -
accuracy: 0.0240 - val_loss: 3.9845 - val_accuracy: 0.0189
Epoch 8/10
239/239 [=====] - 95s 396ms/step - loss: 3.9664 -
accuracy: 0.0239 - val_loss: 3.9713 - val_accuracy: 0.0189
Epoch 9/10
239/239 [=====] - 95s 396ms/step - loss: 3.9637 -
accuracy: 0.0243 - val_loss: 3.9773 - val_accuracy: 0.0189
Epoch 10/10
239/239 [=====] - 96s 399ms/step - loss: 3.9649 -
accuracy: 0.0241 - val_loss: 3.9937 - val_accuracy: 0.0189

```

Summary

- rescale layer
- flatten layer
- dense layer (256 neurons and 38535424 trainable parameters)
- dropout layer
- dense layer (128 neurons and 32896 trainable parameters)
- dropout later
- dense layer (53 neurons and 6837 trainable parameters)

[36]: `model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling_1 (Rescaling)	(None, 224, 224, 3)	0
flatten_1 (Flatten)	(None, 150528)	0
dense_2 (Dense)	(None, 256)	38535424
dropout_2 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 128)	32896
dropout_3 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 53)	6837


=====
 Total params: 38,575,157
 Trainable params: 38,575,157
 Non-trainable params: 0

Plot the Model Metrics for each Epoch

```
[38]: import matplotlib.pyplot as plt

# Plot training & validation accuracy values
plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

score = model.evaluate(test, verbose=0)
print('Test Loss:', score[0])
print('Test Accuracy:', score[1])
```



Test Loss: 3.9936933517456055

Test Accuracy: 0.01886792480945587

Test using an Image

```
[94]: from tensorflow.keras.preprocessing.image import load_img, img_to_array

# Load and preprocess the image
img_path = '/kaggle/input/cards-image-datasetclassification/train/two of clubs/
↪001.jpg'
img = load_img(img_path, target_size=(224, 224))
img_arr = img_to_array(img) / 255.0
img_arr = img_arr.reshape((1, 224, 224, 3))


# Get the model's prediction for the image
prediction = model.predict(img_arr)

# Print the predicted class and its probability
class_index = prediction.argmax()
print(f"Class: {class_names[class_index]}")

# Plot Card Image
plt.imshow(img)
```

```
1/1 [=====] - 0s 43ms/step
Class: two of clubs
```

```
[94]: <matplotlib.image.AxesImage at 0x720b4210e910>
```



Analysis According to the results, the Sequential model achieved a very low test accuracy of approximately 2%. However, it correctly predicted the image class. This could be due to the fact that sequential models are typically designed for analyzing sequential data and may not be the most suitable choice for image analysis. Furthermore, a more complex architecture such as a CNN may be necessary to accurately extract features and classify the images in this dataset.

1.0.5 Convolutional Neural Network Model

- rescale image: normalize pixels values & image size is 224 x 224 and channels = 3
- create a convolutional layer with 32 filters, same padding and relu activation
- max pooling
- create another convolutional layer with 64 filters, same padding and relu activation
- max pooling
- create another convolutional layer with 128 filters, same padding and relu activation
- max pooling
- flattens tensor into 1d array
- drop (0.5) to prevent overfitting
- create a dense layer with 128 neurons and relu activation
- create a dense layer with 53 neurons and softmax activation (used for multi-class classification to output probabilities)

```
[75]: from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dropout, Flatten, Dense

model2 = tf.keras.Sequential([
    tf.keras.layers.Rescaling(1./255, input_shape=(224, 224, 3)),
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(64, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(128, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(num_classes, activation="softmax"),
])
```

Compile/Train Model

```
[76]: # Compile Model
# Loss: categorical_crossentropy (used for multi-class classification for
# labels that must fit in one class)
# Optimizer: adam
# Metrics: Accuracy
```

```

model2.compile(loss='categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

# Train model
# Train model using the training data, 5 epochs, a batch size of 32, and use
↳ the test data
history = model2.fit(train,
                     epochs=10,
                     batch_size=32,
                     verbose=1,
                     validation_data=test)

```

Epoch 1/10

239/239 [=====] - 511s 2s/step - loss: 2.9106 - accuracy: 0.2695 - val_loss: 1.7924 - val_accuracy: 0.4868

Epoch 2/10

239/239 [=====] - 510s 2s/step - loss: 1.7400 - accuracy: 0.5323 - val_loss: 1.4465 - val_accuracy: 0.6264

Epoch 3/10

239/239 [=====] - 508s 2s/step - loss: 1.1108 - accuracy: 0.6929 - val_loss: 1.2129 - val_accuracy: 0.7170

Epoch 4/10

239/239 [=====] - 512s 2s/step - loss: 0.6922 - accuracy: 0.8072 - val_loss: 1.1641 - val_accuracy: 0.7585

Epoch 5/10

239/239 [=====] - 502s 2s/step - loss: 0.4315 - accuracy: 0.8741 - val_loss: 1.2727 - val_accuracy: 0.7660

Epoch 6/10

239/239 [=====] - 506s 2s/step - loss: 0.3062 - accuracy: 0.9138 - val_loss: 1.4705 - val_accuracy: 0.7736

Epoch 7/10

239/239 [=====] - 505s 2s/step - loss: 0.2214 - accuracy: 0.9366 - val_loss: 1.4773 - val_accuracy: 0.7887

Epoch 8/10

239/239 [=====] - 504s 2s/step - loss: 0.1778 - accuracy: 0.9511 - val_loss: 1.4956 - val_accuracy: 0.7811

Epoch 9/10

239/239 [=====] - 505s 2s/step - loss: 0.1640 - accuracy: 0.9557 - val_loss: 1.6589 - val_accuracy: 0.7509

Epoch 10/10

239/239 [=====] - 505s 2s/step - loss: 0.1466 - accuracy: 0.9580 - val_loss: 1.6258 - val_accuracy: 0.7849

Summary

- rescale layer
- 2d conv layer with 32 filters (896 trainable parameters)

- max pooling layer
- 2d conv layer with 64 filters (18496 trainable parameters)
- max pooling layer
- 2d conv layer with 64 filters (73856 trainable parameters)
- max pooling layer
- flatten layer
- dropout layer
- dense layer (128 neurons and 12845184 trainable parameters)
- dense layer (53 neurons and 6837 trainable parameters)

```
[77]: model2.summary()
```


```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
rescaling_2 (Rescaling)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 224, 224, 32)	896
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
conv2d_2 (Conv2D)	(None, 56, 56, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0
flatten_2 (Flatten)	(None, 100352)	0
dropout_4 (Dropout)	(None, 100352)	0
dense_6 (Dense)	(None, 128)	12845184
dense_7 (Dense)	(None, 53)	6837
Total params: 12,945,269		
Trainable params: 12,945,269		
Non-trainable params: 0		

Plot the Model Metrics for each Epoch

```
[78]: plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

score = model2.evaluate(test, verbose=0)
print('Test Loss:', score[0])
print('Test Accuracy:', score[1])
```



```
Test Loss: 1.6257838010787964
Test Accuracy: 0.7849056720733643
```

Test using an Image

```
[107]: # Load and preprocess the image
img_path = '/kaggle/input/cards-image-datasetclassification/train/three of_
↳diamonds/010.jpg'
img = load_img(img_path, target_size=(224, 224))
```

```

img_arr = img_to_array(img) / 255.0
img_arr = img_arr.reshape((1, 224, 224, 3))

# Get the model's prediction for the image
prediction = model2.predict(img_arr)

# Print the predicted class and its probability
class_index = prediction.argmax()
print(f"Class: {class_names[class_index]}")

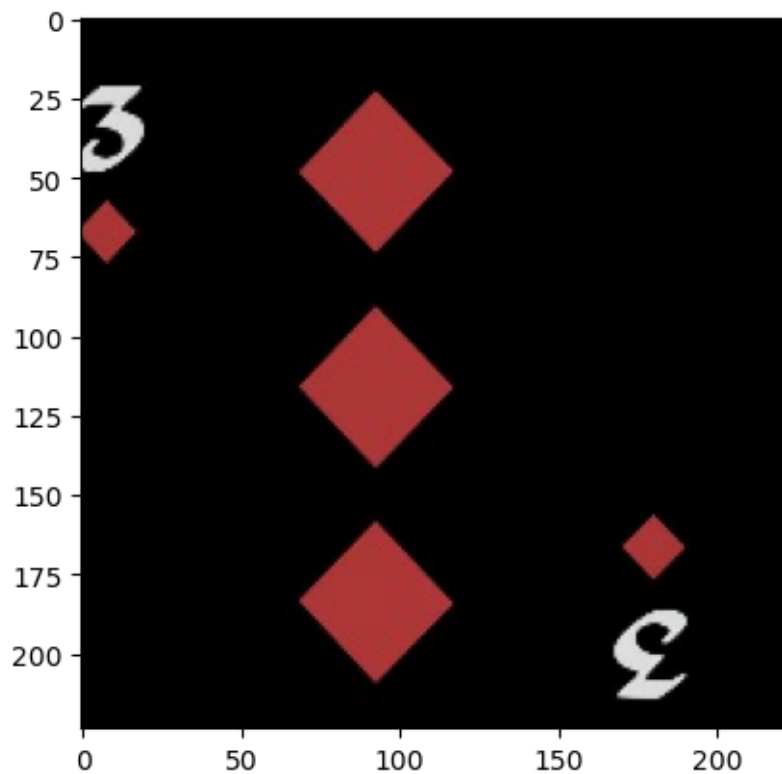
# Plot Card Image
plt.imshow(img)

```

1/1 [=====] - 0s 46ms/step

Class: three of diamonds

[107]: <matplotlib.image.AxesImage at 0x720b40b9ce90>



Analysis Our Convolutional Neural Network (CNN) model has a high accuracy (~78%) and predicted the image correctly due to its ability to learn hierarchical representations, where each layer learns progressively more complex features. Moreover, CNNs reduce the input's dimensionality using convolutional and pooling layers, and control overfitting through methods like dropout.

Additionally, having a large dataset is beneficial to predict better on unseen data.

1.0.6 Pre-Trained Model: MobileNet (V2)

Source: <https://www.kaggle.com/code/paultimothymooney/mobilenetv2-with-tensorflow>

Steps

- Resize the Data: 128 x 128
- Load the Pre-Trained Model
- Define a Sequential Model
- Provide Build Parameters
- Compile the Model
- Train Model

```
[46]: import tensorflow_hub as hub

# Resize the Training Data to 128 x 128
train2 = tf.keras.preprocessing.image_dataset_from_directory(
    "/kaggle/input/cards-image-datasetclassification/train",
    labels="inferred",
    label_mode="categorical",
    color_mode="rgb",
    image_size=(128, 128),
    seed=1234,
    batch_size=32,
)

# Resize the Testing Data to 128 x 128
test2 = tf.keras.preprocessing.image_dataset_from_directory(
    "/kaggle/input/cards-image-datasetclassification/test",
    labels="inferred",
    label_mode="categorical",
    color_mode="rgb",
    image_size=(128, 128),
    seed=1234,
    batch_size=32,
)

# Load the Pre-Trained Model (1st layer)
keras_layer = hub.KerasLayer('https://kaggle.com/models/google/mobilenet-v2/
    ↪frameworks/TensorFlow2/versions/035-128-classification/versions/2')

# Define a Sequential Model:
# Neural Network with 53 neurons (2nd layer)
# Softmax Activation: used for multi-class classification to find probabilities
    ↪for each class (53 classes)
```

```

model3 = tf.keras.Sequential([
    keras_layer,
    tf.keras.layers.Dense(53, activation='softmax')
])

# Build Parameters:
# None: Batch size is not specified
# Shape of the Input: 128 x 128
# Channels: 3
model3.build([None, 128, 128, 3])

# Compile Model
# Loss: categorical_crossentropy (used for multi-class classification for
↳ labels that must fit in one class)
# Optimizer: adam
# Metrics: Accuracy
model3.compile(loss='categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

# Train model
# Train model using the training data, 10 epochs, a batch size of 32, and use
↳ the test data
history = model3.fit(train2,
                    epochs=10,
                    batch_size=32,
                    verbose=1,
                    validation_data=test2)

```

Found 7624 files belonging to 53 classes.

Found 265 files belonging to 53 classes.

Epoch 1/10

239/239 [=====] - 27s 97ms/step - loss: 3.6103 -
accuracy: 0.1177 - val_loss: 3.2053 - val_accuracy: 0.1623

Epoch 2/10

239/239 [=====] - 22s 92ms/step - loss: 3.1194 -
accuracy: 0.1973 - val_loss: 3.0436 - val_accuracy: 0.2075

Epoch 3/10

239/239 [=====] - 23s 94ms/step - loss: 2.9407 -
accuracy: 0.2419 - val_loss: 3.0186 - val_accuracy: 0.2000

Epoch 4/10

239/239 [=====] - 22s 92ms/step - loss: 2.8170 -
accuracy: 0.2638 - val_loss: 3.0124 - val_accuracy: 0.2453

Epoch 5/10

239/239 [=====] - 22s 93ms/step - loss: 2.7505 -
accuracy: 0.2790 - val_loss: 2.9150 - val_accuracy: 0.2377

Epoch 6/10


```

239/239 [=====] - 23s 94ms/step - loss: 2.6671 -
accuracy: 0.3006 - val_loss: 2.8888 - val_accuracy: 0.2415
Epoch 7/10
239/239 [=====] - 23s 95ms/step - loss: 2.6422 -
accuracy: 0.3078 - val_loss: 2.9829 - val_accuracy: 0.2415
Epoch 8/10
239/239 [=====] - 23s 94ms/step - loss: 2.5998 -
accuracy: 0.3169 - val_loss: 2.9271 - val_accuracy: 0.2528
Epoch 9/10
239/239 [=====] - 22s 92ms/step - loss: 2.5372 -
accuracy: 0.3317 - val_loss: 2.9127 - val_accuracy: 0.2453
Epoch 10/10
239/239 [=====] - 22s 93ms/step - loss: 2.5052 -
accuracy: 0.3359 - val_loss: 2.9622 - val_accuracy: 0.2151

```

Summary

- Pre-trained MobileNet (V2) Model (1692489 trainable parameters)
- neural network with 53 classes (53,106 trainable parameters)

```
[47]: model3.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 1001)	1692489
dense_5 (Dense)	(None, 53)	53106

```

=====
Total params: 1,745,595
Trainable params: 53,106
Non-trainable params: 1,692,489
=====

```


Plot the Model Metrics for each Epoch

```

[49]: plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

score = model3.evaluate(test2, verbose=0)
print('Test Loss:', score[0])
print('Test Accuracy:', score[1])

```



Test Loss: 2.962161064147949
 Test Accuracy: 0.21509434282779694

Test using an Image

```
[95]: # Load and preprocess the image
img_path = '/kaggle/input/cards-image-datasetclassification/train/eight of ♣_clubs/011.jpg'
img = load_img(img_path, target_size=(128, 128))
img_arr = img_to_array(img) / 255.0
img_arr = img_arr.reshape((1, 128, 128, 3))


# Get the model's prediction for the image
prediction = model3.predict(img_arr)

# Print the predicted class and its probability
class_index = prediction.argmax()
print(f"Class: {class_names[class_index]}")

# Plot Card Image
plt.imshow(img)
```

```
1/1 [=====] - 0s 30ms/step
Class: queen of clubs
```

```
[95]: <matplotlib.image.AxesImage at 0x720b420fc190>
```



1.0.7 Fine-Tune the Pre-Trained Model: MobileNet (V2)

Source: <https://www.kaggle.com/code/paultimothymooney/mobilenetv2-with-tensorflow>

Steps Define a Sequential Model * Provide Build Parameters * Compile the Model * Train Model

```
[128]: # Define a Sequential Model:
# Rescale
# Neural Network with 128 neurons (2nd layer)
# Dropout to prevent Overfitting
# Neural Network with 53 neurons (3rd layer)
# Softmax Activation: used for multi-class classification to find probabilities
# for each class (53 classes)
model4 = tf.keras.Sequential([
    keras_layer,
    tf.keras.layers.Rescaling(1./127.5, offset=-1),
```

```

        tf.keras.layers.Dense(128, activation='softmax'),
        tf.keras.layers.Dropout(0.5),
        tf.keras.layers.Dense(53, activation='softmax')
    ])

# Build Parameters:
# None: Batch size is not specified
# Shape of the Input: 128 x 128
# Channels: 3
model4.build([None, 128, 128, 3])

# Compile Model
# Loss: categorical_crossentropy (used for multi-class classification for
    ↳ labels that must fit in one class)
# Optimizer: adam (use learning rate)
# Metrics: Accuracy
base_learning_rate = 0.0001
model4.compile(loss='categorical_crossentropy',
               optimizer=tf.keras.optimizers.
    ↳ Adam(learning_rate=base_learning_rate),
               metrics=['accuracy'])

# Train model
# Train model using the training data, 10 epochs, a batch size of 32, and use
    ↳ the test data
history = model4.fit(train2,
                    epochs=10,
                    batch_size=100,
                    verbose=1,
                    validation_data=test2)

```

```

Epoch 1/10
239/239 [=====] - 25s 98ms/step - loss: 3.9704 -
accuracy: 0.0192 - val_loss: 3.9732 - val_accuracy: 0.0189
Epoch 2/10
239/239 [=====] - 23s 96ms/step - loss: 3.9698 -
accuracy: 0.0212 - val_loss: 3.9709 - val_accuracy: 0.0189
Epoch 3/10
239/239 [=====] - 23s 95ms/step - loss: 3.9694 -
accuracy: 0.0219 - val_loss: 3.9710 - val_accuracy: 0.0189
Epoch 4/10
239/239 [=====] - 22s 93ms/step - loss: 3.9681 -
accuracy: 0.0202 - val_loss: 3.9711 - val_accuracy: 0.0189
Epoch 5/10
239/239 [=====] - 23s 95ms/step - loss: 3.9682 -
accuracy: 0.0210 - val_loss: 3.9712 - val_accuracy: 0.0189
Epoch 6/10

```

```

239/239 [=====] - 23s 95ms/step - loss: 3.9678 -
accuracy: 0.0222 - val_loss: 3.9717 - val_accuracy: 0.0189
Epoch 7/10
239/239 [=====] - 23s 96ms/step - loss: 3.9688 -
accuracy: 0.0209 - val_loss: 3.9715 - val_accuracy: 0.0189
Epoch 8/10
239/239 [=====] - 23s 95ms/step - loss: 3.9672 -
accuracy: 0.0232 - val_loss: 3.9713 - val_accuracy: 0.0189
Epoch 9/10
239/239 [=====] - 23s 95ms/step - loss: 3.9669 -
accuracy: 0.0249 - val_loss: 3.9714 - val_accuracy: 0.0189
Epoch 10/10
239/239 [=====] - 23s 95ms/step - loss: 3.9670 -
accuracy: 0.0219 - val_loss: 3.9718 - val_accuracy: 0.0189

```

Summary

- Pre-trained MobileNet (V2) Model (1692489 trainable parameters)
- rescale layer
- dense layer (128 neurons and 128256 trainable parameters)
- dropout layer
- dense layer (53 neurons and 6837 trainable parameters)

[129]: `model14.summary()`

Model: "sequential_18"

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 1001)	1692489
rescaling_8 (Rescaling)	(None, 1001)	0
dense_25 (Dense)	(None, 128)	128256
dropout_10 (Dropout)	(None, 128)	0
dense_26 (Dense)	(None, 53)	6837

```


=====
Total params: 1,827,582
Trainable params: 135,093
Non-trainable params: 1,692,489
-----

```

Plot the Model Metrics for each Epoch

```
[130]: plt.plot(history.history['val_accuracy'])
plt.plot(history.history['accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')
plt.show()

score = model4.evaluate(test2, verbose=0)
print('Test Loss:', score[0])
print('Test Accuracy:', score[1])
```



Test Loss: 3.971834897994995
Test Accuracy: 0.01886792480945587

Test using an Image

```
[131]: # Load and preprocess the image
img_path = '/kaggle/input/cards-image-datasetclassification/train/eight of_
clubs/011.jpg'
img = load_img(img_path, target_size=(128, 128))
```

```

img_arr = img_to_array(img) / 255.0
img_arr = img_arr.reshape((1, 128, 128, 3))

# Get the model's prediction for the image
prediction = model3.predict(img_arr)

# Print the predicted class and its probability
class_index = prediction.argmax()
print(f"Class: {class_names[class_index]}")

# Plot Card Image
plt.imshow(img)


```

```

1/1 [=====] - 0s 30ms/step
Class: queen of clubs

```

```
[131]: <matplotlib.image.AxesImage at 0x720b282edf90>
```



Analysis As observed, the test accuracy obtained through transfer learning using the MobileNet (V2) model is quite low (~22% & ~2%). Additionally, the model predictions were wrong both times. One possible explanation for this could be that the MobileNet (V2) model was originally trained using the massive ImageNet dataset and may not be able to generalize well on our dataset, which is

much smaller in comparison. Specifically, the ImageNet dataset contains 14,197,122 images, whereas ours only has 7,624 images. Hence, it is likely that the MobileNet (V2) model is too complex to effectively learn from our limited amount of data. It is worth noting that the pre-trained model used in this case is the one from Kaggle, not the original MobileNet V2 model developed by Google. This could also potentially explain the low accuracy observed in the results.

1.1 The CNN Model yields the highest accuracy for our data!